

Spinach: An Ad-hoc Query Engine on Top of Spark SQL

Cheng Hao(hao.cheng@intel.com)

Outlines

- Major features
- Getting Starts
- Implementation & Design principles
- Micro benchmark the demo data
- Future Plan

About Me

- Active Spark Contributor in Apache Open Source
- Engineering Manager from BDT of Intel APAC
- Leading the IA Optimization for Spark at Intel


How to accelerate SQL queries with Spark SQL?

- Tungsten
 - (Offheap) Data oriented Memory Management
 - Cache-aware computation
 - Code Generation
- Tungsten II
 - Whole Stage Code Generation
 - Vectorization
- ...

What Else?

Spark

Computing Engine



Data Source API

ALLUXIO

redis

cassandra

APACHE HBASE

Cache Layer

hadoop HDFS

amazon web services S3

阿里云 OSS
aliyun.com

Storage Layer

Spark

Computing Engine

No additional 3rd Service required

Fine-grained Data Cached

Data Source API

Spinach

Customized Indices Supported

Cache Layer

Data Cached in Off-heap Memory(No GC Overhead)



Storage Layer

Getting Started

1. Start the Spark SQL Shell and Load the Spinach Package

```
$SPARK_HOME/bin/spark-sql --jars spinach-0.1.jar
```

2. Create a Spinach backend Data Source Table

```
spark-sql> CREATE TABLE src(a INT, b STRING, value INT) USING org.apache.spark.sql.execution.datasources.spinach;
```

3. Add Index Support the Data Source Table

```
spark-sql>CREATE INDEX idx_1 ON src (a);
```

4. Ad-hoc Query by auto enable the indices

```
spark-sql> INSERT INTO TABLE src SELECT key1, key2, value FROM xxx;
```

```
spark-sql> SELECT MAX(value) FROM src WHERE a > 100 AND a <= 120 AND b='spinach';
```

```
spark-sql> CREATE INDEX idx_2 ON src (a, b);
```

```
spark-sql> SELECT MAX(value) FROM src WHERE a>=100 AND b='spinach';
```

```
spark-sql> DROP INDEX idx_2;
```

```
spark-sql> SELECT MAX(value) FROM src WHERE a>=100;
```

Auto trigger the index idx_1



Auto trigger the index idx_2(TBD)

Trigger the index idx_1, but found too many records return, auto bypass index and fall back to full table scan

- **DDL Statement Extension (Index Management)**
 - Create / Add Index (Parser & Logical Node / Physical Execution)
 - Drop Index (Parser & Logical Node / Physical Execution)
- **Data Source Extension**
 - Implements the HadoopFSRelation interface (Support Partition & File Status Caching)
 - Abbr. (“**spn**”) for Spinach Data Source in Data Frame API

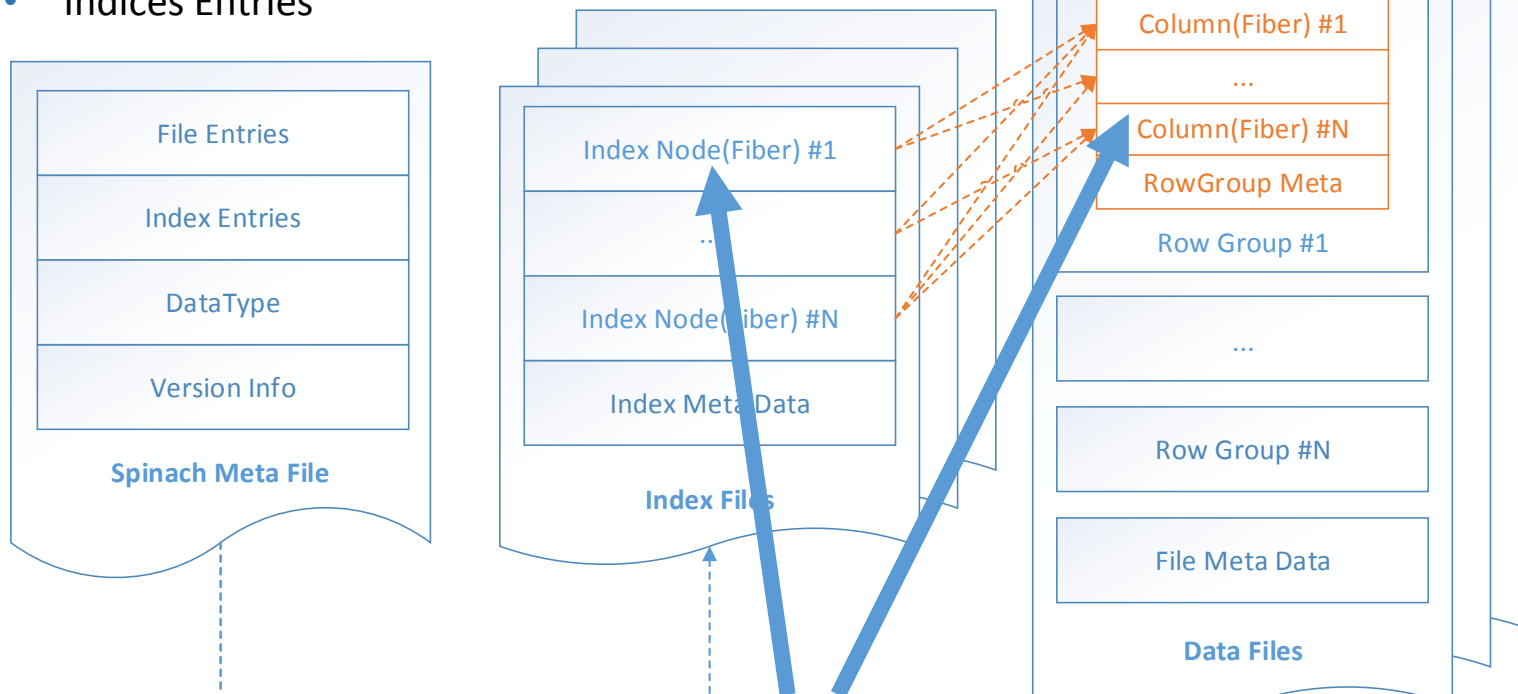
```
df.write.format("spn").save("/path/to/spinach_test")  
sqlContext.read.format("spn").load("/path/to/spinach_test")
```

- **Enable the extensions**
 - SpinachContext (SQLContext)
 - Make SQLContext configurable in ThriftServer / SparkSQL Shell / Spark-shell
 - Spark Executor HeartBeat extensions (talk later)

- Data File (N Files)
 - RowGroups
 - Fibers in Each Row Group
 - File Meta


- Index File (N * M Files)
 - N is the number of Data Files
 - M is the number of Indices
 - Index Meta
 - Index Fibers

- Spinach Meta (1)
 - Data Schema
 - Data File Statistic / Entries
 - Indices Entries




- Fibers (the minimum unit for caching / loading / eviction)
 - Index Fibers
 - Data Fibers (Columnar based)

INDEX of UID



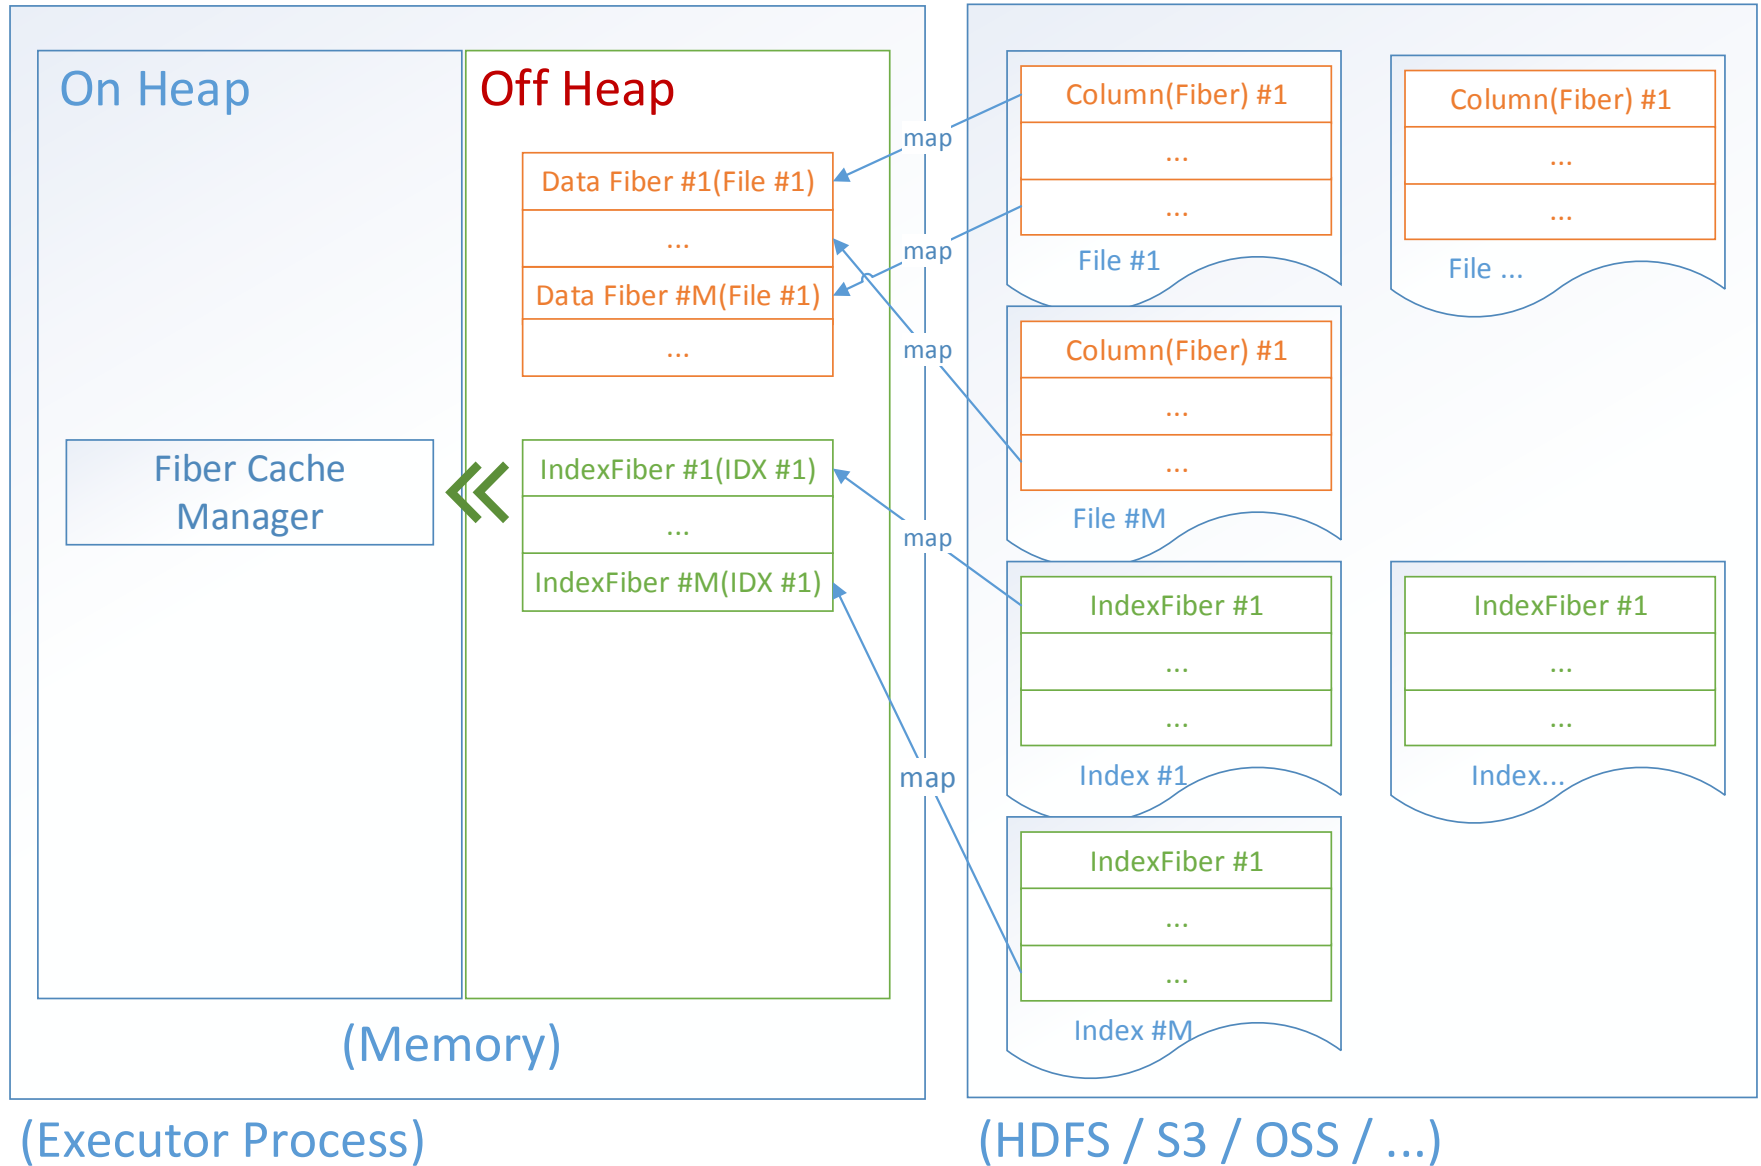
INDEX of Age



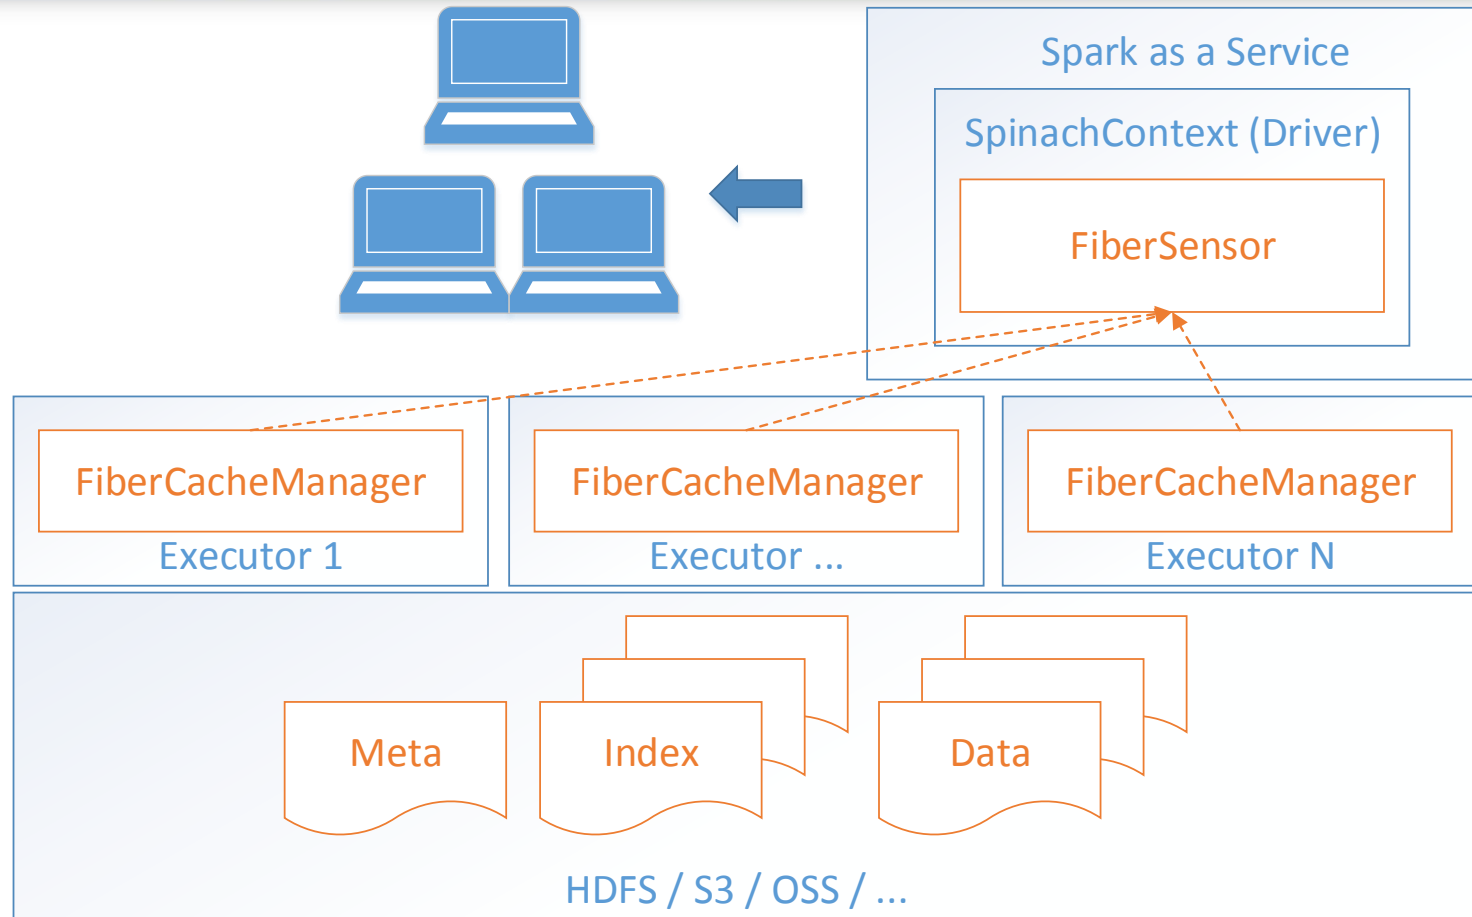
Row Group

UID	Age	Name	
2	18	二狗子	Row 1
5	25	王五	Row 2
4	53	李四	Row 3
3	34	张三	Row 4
6	17	赵六	Row 5
7	18	钱多多	Row 6


Fibers are in the red boxes



- Spinach Meta
 - Describe the data schema and statistic info
 - Describe how the indices are organized
 - Different cache strategy for fast accessing
- Data Fiber
 - Columnar Storage
 - Aim to fully compatible with Spark SQL Data Types (nested data types are TBD)
 - Vectorization friendly in the offheap Memory, without any encoding
 - Data Type aware encoder/decoder in the storage layer (TBD)
 - Decouple with concrete data format to support more columnar storage based format like ORC, Parquet (TBD)
- Index Fiber (Sort based)
 - Row(index keys) based Storage
 - MySQL like B+ Tree Implementation
 - Separate the files for index & data, for better managing the indices efficiently, and decouple with the data format.



- Fiber Cache Manager
 - Resides in each executor process
 - Manage an off-heap memory pool & Data Loading & Evicting Strategy
 - Update the fiber cache statistic info periodically with Fiber Sensor via executor heartbeat RPC.



- Fiber Sensor
 - Resides in Spark Driver Process
 - Global Fiber Cache Distribution statistic info
 - Fiber(index/data) cache-aware for preferred data location in tasks assignment

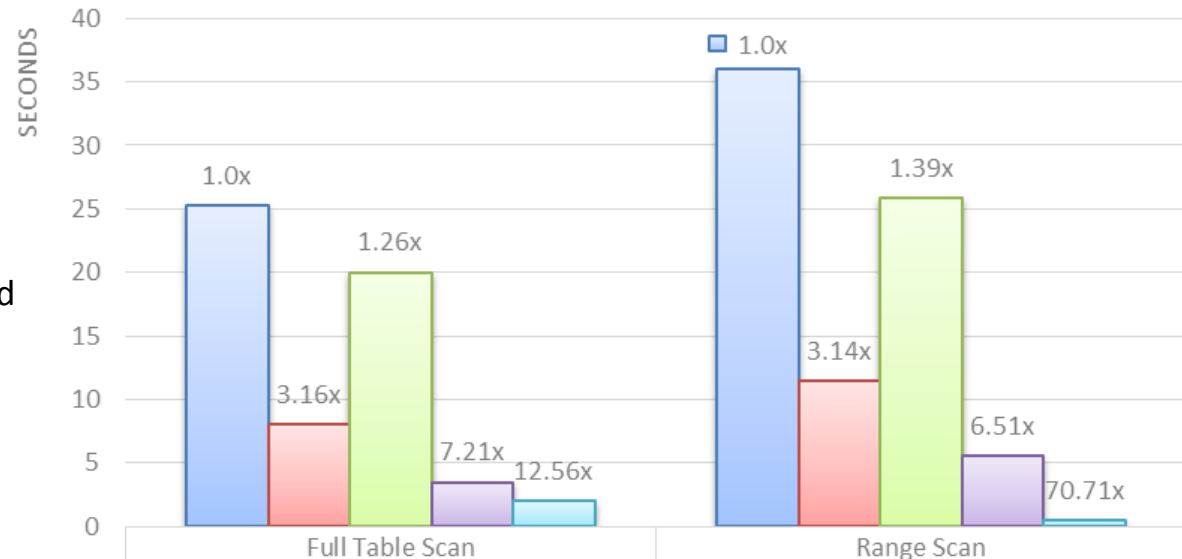
Hardware

- 1 Master + 3 Slaves
- Xeon E5-2699 v3
- 256GB / node
- 8x1TB SATA / node
- 10Gb network

Data

- ~100GB Uncompressed
- Record Count: 1.3B

Demo Micro Benchmark



■ Parquet(Compressed Cold Data)	25.25	36.06
■ Parquet(Compressed) (OS Cached)	8	11.5
■ Parquet(NonCompression) (OS Cached)	20	25.9
■ Spark InMemory Cache*	3.5	5.54
■ Spinach	2.01	0.51

- Demo Micro-benchmark probably very different when data value patterns are different*
- Full table scan:
 - `df.selectExpr("count(str1)", "count(int1)", "count(str2)").show`
- Range Key Scan: ()
 - `df.filter("str2 >= 'China-6234567' and str2 <= 'China-6234596']").selectExpr("count(str1)", "sum(int1)").show`

- Data Source Extension
- Fined-grained Data Cache
- User Defined Indices
- Vectorization Friendly
- Off-heap Memory

- Open Source (POC stage)
- More Index type (e.g. bloom filter) & better encoder/decoder
- Nested data type
- Support other columnar storage based data formats (ORC/Parquet)
- More Flexible Fiber Caching & Evicting strategy
- Optimize task assignment algorithm(preferred location)
- Auto-detect & cache the shared common sub queries result

Your Idea Matters!!!

Thanks